

Assignment 2

Sheil Sarda
sheils@seas.upenn.edu

September 20, 2020

Exercise 1

(a) Scale: down-samples the input array by a factor of two, transforming an image of size $HEIGHT * WIDTH$ to an image of size $(HEIGHT/2) * (WIDTH/2)$. The index used inside the nested for-loop ($Y * WIDTH + X$) is used to access the pixel located at (X, Y) , assuming the origin of the coordinate system is located in the top-left corner.

```
void Scale(const unsigned char * Input, unsigned char * Output){
    for (int Y = 0; Y < HEIGHT; Y += 2)
        for (int X = 0; X < WIDTH; X += 2)
            Output[(Y / 2) * WIDTH / 2 + (X / 2)] = Input[Y * WIDTH + X];
}
```

(b) Filter:

- Filter applies a horizontal and vertical filter to the input image, and returns the filtered image as the output.
- Filter_horizontal scans through every pixel of the image, and replaces it with SUM divided by 256. SUM is the value of each Coefficient, where coefficients add up to 256, multiplied by the value stored in each pixel. In effect, this smooths each pixel of the image by replacing the value with a weighted average of the pixels in the same row, or for pixels on the edge– a weighted average of pixels in the adjacent row.
- Filter_vertical functions the same way as Filter_horizontal except that it replaces each pixel with a weighted average of the pixels in the same column, or for pixels on the edge– a weighted average of pixels in adjacent columns.

```
void Filter(const unsigned char * Input, unsigned char * Output){
    unsigned char * Temp = (unsigned char*)malloc(INPUT_HEIGHT * OUTPUT_WIDTH);
    Filter_horizontal(Input, Temp);
    Filter_vertical(Temp, Output);
    free(Temp);
}
```

```
void Filter_horizontal(const unsigned char * Input, unsigned char * Output){
    for (int Y = 0; Y < INPUT_HEIGHT; Y++){
        for (int X = 0; X < OUTPUT_WIDTH; X++){
            unsigned int Sum = 0;
            for (int i = 0; i < FILTER_LENGTH; i++){
                Sum += Coefficients[i] * Input[Y * INPUT_WIDTH + X + i];
            }
            Output[Y * OUTPUT_WIDTH + X] = Sum >> 8;
        }
    }
}
```

```
static void Filter_vertical(const unsigned char * Input, unsigned char * Output){
    for (int Y = 0; Y < OUTPUT_HEIGHT; Y++){
        for (int X = 0; X < OUTPUT_WIDTH; X++){
            unsigned int Sum = 0;
            for (int i = 0; i < FILTER_LENGTH; i++){
                Sum += Coefficients[i] * Input[(Y + i) * OUTPUT_WIDTH + X];
            }
            Output[Y * OUTPUT_WIDTH + X] = Sum >> 8;
        }
    }
}
```

(c) Differentiate: Replaces each pixel of the image with the difference of the input and AVERAGE. AVERAGE is defined as the arithmetic mean of the value of the pixels located directly above and to the left of the current pixel. For edge cases where the previous row or column may not exist, we simply use either previous row or previous column.

```
void Differentiate(const unsigned char * Input, unsigned char * Output){
    for (int Y = 0; Y < HEIGHT; Y++){
        for (int X = 0; X < WIDTH; X++){
            int Average = 0;
            if (Y > 0 && X > 0)
                Average = (Input[WIDTH * (Y - 1) + X] + Input[WIDTH * Y + X - 1]) / 2;
            else if (Y > 0)
                Average = Input[WIDTH * (Y - 1) + X];
            else if (X > 0)
                Average = Input[WIDTH * Y + X - 1];
            unsigned char Diff = Input[WIDTH * Y + X] - Average;
            Output[Y * WIDTH + X] = Diff;
        }
    }
}
```

(d) Compress: for each pixel in the image, first fetch the corresponding Code and Code_length using the pixel value. Then, computer Byte by bit-shifting the current value to the right, and filling it in with the bit-sequence in Code, but reverse order. When Byte reaches a length of 8 digits, it gets put into the Output array, and the local variable is reset to 0 to start again. The Output array is thus $\frac{1}{8}$ the size of the Input array, and each value in the array is a Code.

```
int Compress(const unsigned char * Input, unsigned char * Output){
    unsigned int Byte = 0;
    int Length = 0;
    for (int i = 0; i < SIZE; i++){
        unsigned long long Code = Codes[Input[i]];
        int Code_length = Code_lengths[Input[i]];

        for (int j = 0; j < Code_length; j++){
            Byte = (Byte << 1) | ((Code >> (Code_length - 1 - j)) & 1);

            if (++Length % 8 == 0){
                Output[Length / 8 - 1] = Byte;
                Byte = 0;
            }
        }
    }
    if (Length % 8 > 0)
```

```

    Output[Length / 8] = Byte;
    return Length / 8 + 1;
}

```

Exercise 2

(a)

Functions	Average Latency ($T_{measured_avg}$ ns)	% of total latency (1.28811e+08 ns)	Average Latency (2.3 GHz) ($T_{measured_avg}$ cycles)	Actual Latency (% of total cycles)
Scale	3.93261e+06	3.05%	9,045,003	9,237,036
Filter horizontal	2.94201e+07	22.84%	67,666,230	69,171,771
Filter vertical	3.26903e+07	25.38%	75,187,690	76,864,253
Differentiate	8.56184e+06	6.65%	19,692,232	20,139,767
Compress	4.59722e+07	35.69%	105,736,060	108,087,414

Exercise 3

(a) From the above table, the function with the highest latency alone is compress.

If we include filter in this comparison, which is comprised of filter horizontal and vertical, it would have the highest latency.

(b) DFG of the body of the loop over x :

Cycle	Multiply (M)	Accumulate (A)	Shift (S)
S1	M1 = Coeff[0] * Input[0]		
S2	M2 = Coeff[1] * Input[0]	Sum += M1	
S3	M3 = Coeff[2] * Input[0]	Sum += M2	
S4	M4 = Coeff[3] * Input[0]	Sum += M3	
...			
S Filter Length + 2			Output[0] = Sum >> 8
S Filter Length + 3	M1 = Coeff[0] * Input[1]		
...			

Figure 1: Instructions executed each cycle, excluding memory accesses for INPUT and Coefficients array, since those instructions were not mentioned in the submission guidelines.

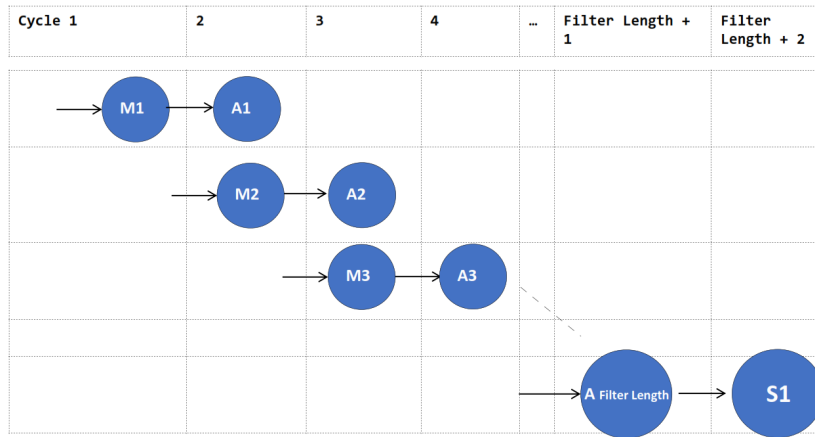


Figure 2: Data Flow Graph with timing information

```

void Filter_horizontal(const unsigned char * Input, unsigned char * Output) {
    for (int Y = 0; Y < INPUT_HEIGHT*OUTPUT_WIDTH; Y++) {
        unsigned int Sum = 0;
        for (int i = 0; i < FILTER_LENGTH; i++)
            Sum += Coefficients[i] * Input[Y + i];

        Output[Y] = Sum >> 8;
    }
}

```

(c) Total number of compute operations equals $(\text{FILTER_LENGTH} * 3 + 2) * (1500 * 2000) = 69,000,000$. This number excludes loop overhead, and assumes that memory fetches from the INPUT and Coefficients array can occur in one cycle each.

(d) The $2\times$ speed-up should be applied to the slowest segment of the code, which as we discussed above, is the compress function. This function currently comprises $\sim 36\%$ of the total latency.

(e) Assuming we speedup the compress function by a factor of $2\times$, $T_{after} = \frac{1}{2} * 35.69\% + 1 * (1 - 35.69\%) = 0.82$. The resulting speedup = $\frac{T_{before}}{T_{after}} = \frac{1}{0.82} = 1.22\times$.




Figure 3: Data Flow Graph assuming unlimited resources.

(f) New Data Flow Graph with lowest critical path delay.

(g) New critical path length assuming instructions can execute in the same cycle in terms of compute operations is $\text{FILTER_LENGTH} + 2$ cycles assuming addition, multiplication and bit-shift operations take a single cycle to execute.

(h) Resource capacity lower bound for the loop body assuming 4 multipliers, 2 adders and 1 shifter is $\frac{\text{TotalOps}}{\text{Operators}}$

- Total # of add operations: $7 * (1500 * 2000) = 21,000,000$. Since we have 2 adders, $RB_{ADD} = 10,500,000$.
- Total # of multiplication operations is 21,000,000 using the same logic. Since we have 4 multipliers, $RB_{MULT} = 5,250,000$.
- Total # of shift operations is 3,000,000. Since we have 1 shifter, $RB_{SHIFT} = 3,000,000$.

The overall Resource Bound of the algorithm is the max of the above bounds. In this case, $RB = 10,500,000$ cycles.

For a single loop iteration,

- $RB_{ADD} = 3.5$.
- $RB_{MULT} = 1.75$.
- $RB_{SHIFT} = 1$.

The RB would equal 3.5 cycles, bottle-necked by the adders.

Exercise 4

(a)

```
(gdb) n
60          Sum += Coefficients[i] * Input[Y * INPUT_WIDTH + X + i];
(gdb) disassemble
Dump of assembler code for function Filter_horizontal:
0x00000000004014a0 <+0>:      add     x9, x0, \#0x2dc, lsl \#12
0x00000000004014a4 <+4>:      adrp    x6, 0x420000
0x00000000004014a8 <+8>:      add     x8, x0, \#0x7ca
0x00000000004014ac <+12>:     add     x9, x9, \#0xe8a
=> 0x00000000004014b0 <+16>:     add     x6, x6, \#0xcc0
0x00000000004014b4 <+20>:     sub     x5, x8, \#0x7ca
0x00000000004014b8 <+24>:     mov     x7, x1
0x00000000004014bc <+28>:     nop
0x00000000004014c0 <+32>:     mov     x0, \#0x0                                // \#0
0x00000000004014c4 <+36>:     mov     w2, \#0x0                                // \#0
0x00000000004014c8 <+40>:     ldrb    w4, [x5, x0]
0x00000000004014cc <+44>:     ldr     w3, [x6, x0, lsl \#2]
0x00000000004014d0 <+48>:     add     x0, x0, \#0x1
0x00000000004014d4 <+52>:     cmp     x0, \#0x7
0x00000000004014d8 <+56>:     madd    w2, w4, w3, w2
0x00000000004014dc <+60>:     b.ne    0x4014c8 <Filter_horizontal+40> // b.any
0x00000000004014e0 <+64>:     lsr     w2, w2, \#8
0x00000000004014e4 <+68>:     strb    w2, [x7], \#1
0x00000000004014e8 <+72>:     add     x5, x5, \#0x1
0x00000000004014ec <+76>:     cmp     x8, x5
0x00000000004014f0 <+80>:     b.ne    0x4014c0 <Filter_horizontal+32> // b.any
0x00000000004014f4 <+84>:     add     x8, x8, \#0x7d0
0x00000000004014f8 <+88>:     add     x1, x1, \#0x7ca
0x00000000004014fc <+92>:     cmp     x9, x8
0x0000000000401500 <+96>:     b.ne    0x4014b4 <Filter_horizontal+20> // b.any
0x0000000000401504 <+100>:    ret
End of assembler dump.
```

(b)

Assembly Instructions	Annotation	# of function calls (N)	# of cycles per call (T)	# of instructions issued (N_{issue})	Total cycles ($T \times N$)
add x6, x6, #0xcc0	loop incr	1500	1	3	500
sub x5, x8, #0x7ca	fetch var	1500	1	3	500
mov x7, x1	fetch var	2000	1	3	667
mov x0, #0x0	reset sum	3,000,000	1	3	1,000,000
mov x0, #0x0	reset sum	3,000,000	1	3	1,000,000
ldrb w4, [x5, x0]	array read	21,000,000	1	3	7,000,000
ldr w3, [x6, x0, lsl #2]	array read	21,000,000	1	3	7,000,000
add x0, x0, #0x1	index add	21,000,000	1	3	7,000,000
cmp x0, #0x7	loop limit	21,000,000	1	3	7,000,000
madd w2, w4, w3, w2	coeff mult	21,000,000	1	3	7,000,000
b.ne 0x4014c8	branch	21,000,000	1	3	7,000,000
lsr w2, w2, #8	div by 256	3,000,000	1	3	1,000,000
strb w2, [x7], #1	store out	3,000,000	1	3	1,000,000
add x5, x5, #0x1	index add	3,000,000	1	3	1,000,000
cmp x8, x5	loop limit	3,000,000	1	3	1,000,000
b.ne 0x4014c0	branch	3,000,000	1	3	1,000,000
add x8, x8, #0x7d0	add to sum	1500	1	3	500
add x1, x1, #0x7ca	add to sum	1500	1	3	500
cmp x9, x8	loop limit	1500	1	3	500
b.ne 0x4014b4	branch	1500	1	3	500
ret	return	1	1	3	0
					49,003,667

(c) Other instructions in the above table are used to:

- Divide Sum by 256 before array store. This is performed through bit-shifting operations.
- Store new pixel value to the Output array.
- Fetch variables to a register for loop limit comparisons.
- Reset Sum to 0 at the beginning of the iteration.
- Return statement at the end of the function.

(d) Completed in table.

(e) Compute latency of non-memory instructions.

Assembly Instructions	Annotation	# of function calls (N)	# of cycles per call (T)	# of instructions issued (N_{issue})	Total cycles ($T \times N$)
add x6, x6, #0xcc0	loop incr	1500	1	3	500
sub x5, x8, #0x7ca	fetch var	1500	1	3	500
mov x7, x1	fetch var	2000	1	3	667
mov x0, #0x0	reset sum	3,000,000	1	3	1,000,000
mov x0, #0x0	reset sum	3,000,000	1	3	1,000,000
add x0, x0, #0x1	index add	21,000,000	1	3	7,000,000
cmp x0, #0x7	loop limit	21,000,000	1	3	7,000,000
madd w2, w4, w3, w2	coeff mult	21,000,000	1	3	7,000,000
b.ne 0x4014c8	branch	21,000,000	1	3	7,000,000
lsr w2, w2, #8	div by 256	3,000,000	1	3	1,000,000
add x5, x5, #0x1	index add	3,000,000	1	3	1,000,000
cmp x8, x5	loop limit	3,000,000	1	3	1,000,000
b.ne 0x4014c0	branch	3,000,000	1	3	1,000,000
add x8, x8, #0x7d0	add to sum	1500	1	3	500
add x1, x1, #0x7ca	add to sum	1500	1	3	500
cmp x9, x8	loop limit	1500	1	3	500
b.ne 0x4014b4	branch	1500	1	3	500
ret	return	1	1	3	0
					34,003,667

$$T_{cycle_memory} = T_{total} - T_{non_memory} = 67,666,230 - 34,003,667 = 33,662,563$$

(f) Identified memory operations, and if the ops are to memory locations not loaded during invocation of the function.

Assembly Instructions	Annotation	# of function calls (N)	Previously observed (Y/N)	% of total cycles
ldrb w4, [x5, x0]	array read	21,000,000	Y (coeff.)	14.29%
ldr w3, [x6, x0, lsl #2]	array read	21,000,000	N (Input)	14.29%
strb w2, [x7], #1	output	3,000,000	N (Output)	2.04%

Estimates for different quantities in the timing equation:

$$N_{total} = 147,011,001 \text{ cycles}$$

$$N_{slow_loads} = 24,000,000 \text{ cycles}$$

$$N_{fast_loads} = 21,000,000 \text{ cycles}$$

(g) We know:

$$T_{cycle_memory} = 33,662,563 = T_{fast_loads} + T_{slow_loads}$$

Also:

$$T_{fast_loads} = \frac{21,000,000}{3} \times 1 = 7,000,000$$

Substituting in the original equation, we get:

$$T_{slow_loads} = 33,662,563 - 7,000,000 = 26,662,563 \text{ cycles}$$